

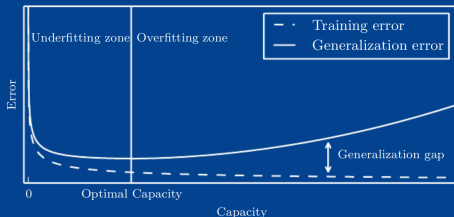


## LESSON 7: Koncepts II

### capacity, under- and overfit, generalization

CARSTEN EIE FRIGAARD

AUTUMN 2019



# Supergruppe diskussion

Diskussion ml. Grp A og B:

- ▶ ca. 15 min: Grp A genfortæller første halvdel af § 2
  - ▶ "Look at the Big Picture",
  - ▶ "Get the Data",  
    eksklusiv "Create the Workspace" og "Download the Data",
  - ▶ "Discover and Visualize the Data to Gain Insights",
  - ▶ "Prepare the Data for Machine Learning Algorithms",
  - ▶ "Select and Train a Model",
- ▶ ca 15 min: herefter byttes og Grp B genfortæller..
  - ▶ "Fine-Tune Your Model",
  - ▶ "Launch, Monitor, and Maintain Your System",
  - ▶ "Try It Out!".

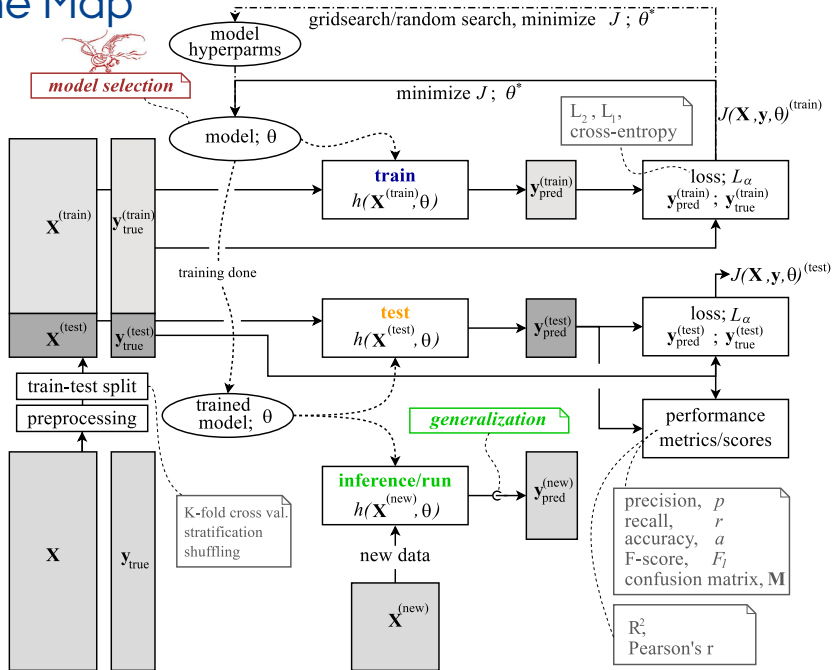
Inddrag og relater til jeres eget **slut-projekt**,

Inddrag **The Map** for Supervised learning

Session I: 8:30..9:15

Session II: 9:15..10:00

# The Map



# Pipelines

Brief intro to Scikit-learn pipelines..

Python code from capacity\_under\_overfitting.ipynb

```
1  from sklearn.pipeline import Pipeline
2  from sklearn.preprocessing import PolynomialFeatures
3  from sklearn.linear_model import LinearRegression
4  from sklearn.model_selection import cross_val_score
5
6  ..
7
8  polynomial_features = PolynomialFeatures(degree=degrees[i], ..
9  linear_regression = LinearRegression()
10
11  pipeline = Pipeline([
12      ("polynomial_features", polynomial_features),
13      ("linear_regression", linear_regression)
14  ])
15  pipeline.fit(X[:, np.newaxis], y)
16
17  scores = cross_val_score(pipeline, X[:, np.newaxis], y, scoring=
18      "neg_mean_squared_error", cv=10)
19  ..
20  score_mean = -scores.mean()
```

# RESUMÉ: L02/performance\_metrics.ipynb

## Classification metrics


See the [Classification metrics](#) section of the user guide for further details.

<code>metrics.accuracy_score(y_true, y_pred[, ...])</code>	Accuracy classification score.
<code>metrics.auc(x, y[, reorder])</code>	Compute Area Under the Curve (AUC) using the trapezoidal rule
<code>metrics.average_precision_score(y_true, y_score)</code>	Compute average precision (AP) from prediction scores
<code>metrics.cohen_kappa_score(y1, y2[, labels, ...])</code>	Cohen's kappa: a statistic that measures inter-annotator agreement.
<code>metrics.confusion_matrix(y_true, y_pred[, ...])</code>	Compute confusion matrix to evaluate the accuracy of a classification
<code>metrics.f1_score(y_true, y_pred[, labels, ...])</code>	Compute the F1 score, also known as balanced F-score or F-measure
<code>metrics.log_loss(y_true, y_pred[, eps, ...])</code>	Log loss, aka logistic loss or cross-entropy loss.
<code>metrics.precision_score(y_true, y_pred[, ...])</code>	Compute the precision
<code>metrics.recall_score(y_true, y_pred[, ...])</code>	Compute the recall
<code>metrics.roc_auc_score(y_true, y_score[, ...])</code>	Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores.
<code>metrics.roc_curve(y_true, y_score[, ...])</code>	Compute Receiver operating characteristic (ROC)
<code>metrics.zero_one_loss(y_true, y_pred[, ...])</code>	Zero-one classification loss.

## Regression metrics

See the [Regression metrics](#) section of the user guide for further details.

<code>metrics.explained_variance_score(y_true, y_pred)</code>	Explained variance regression score function
<code>metrics.max_error(y_true, y_pred)</code>	max_error metric calculates the maximum residual error.
<code>metrics.mean_absolute_error(y_true, y_pred)</code>	Mean absolute error regression loss
<code>metrics.mean_squared_error(y_true, y_pred[, ...])</code>	Mean squared error regression loss
<code>metrics.mean_squared_log_error(y_true, y_pred)</code>	Mean squared logarithmic error regression loss
<code>metrics.median_absolute_error(y_true, y_pred)</code>	Median absolute error regression loss
<code>metrics.r2_score(y_true, y_pred[, ...])</code>	R <sup>2</sup> (coefficient of determination) regression



# Model capacity

Exercise: `capacity_under_overfitting.ipynb`

Dummy and Paradox classifier:

*capacity* fixed  $\sim 0$ , cannot generalize at all!

Linear regression for a polynomial model:

*capacity*  $\sim$  degree of the polynomial,  $x^n$

Neural Network model:

*capacity*  $\propto$  number of neurons/layers


$\Rightarrow$  **Capacity** can be hard to express as a quantity for some models, but you need to choose..

$\Rightarrow$  how to choose the **optimal capacity**?

# Under- and overfitting

Exercise: `capacity_under_overfitting.ipynb`

Polynomial linear reg. fit for underlying model:  $\cos(x)$



- ▶ underfitting: capacity of model too low,
- ▶ overfitting: capacity too high.


⇒ how to choose the **optimal** capacity?

# Generalization Error


Exercise: `generalization_error.ipynb`

RMSE-capacity plot for lin. reg. with polynomial features

(capacity  $\sim$  degree of poly)



(Figure 5.3 from [DL])



Inspecting the plots from the exercise (`.ipynb`) and [DL],  
extracting the concepts:

- ▶ training/generalization error,
- ▶ generalization gap,
- ▶ underfit/overfit zone,
- ▶ optimal capacity (best-model, early stop),
- ▶ (and the two axes: x/capacity, y/error.)



# Generalization Error

Exercise: `generalization_error.ipynb`

**NOTE:** three methods/plots:

- i) via **learning curves** as in [HOML],
- ii) via an **error-capacity** plot as in [GITHOLM] and [DL],
- ii) via an **error-epoch** plot as in [GITHOML].

