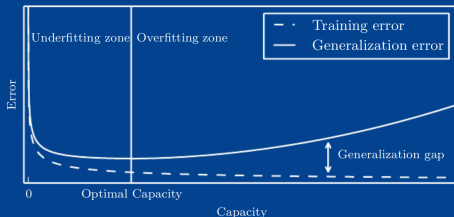




## LESSON 8: Model-capacity, Under- and Overfitting, Generalization

CARSTEN EIE FRIGAARD

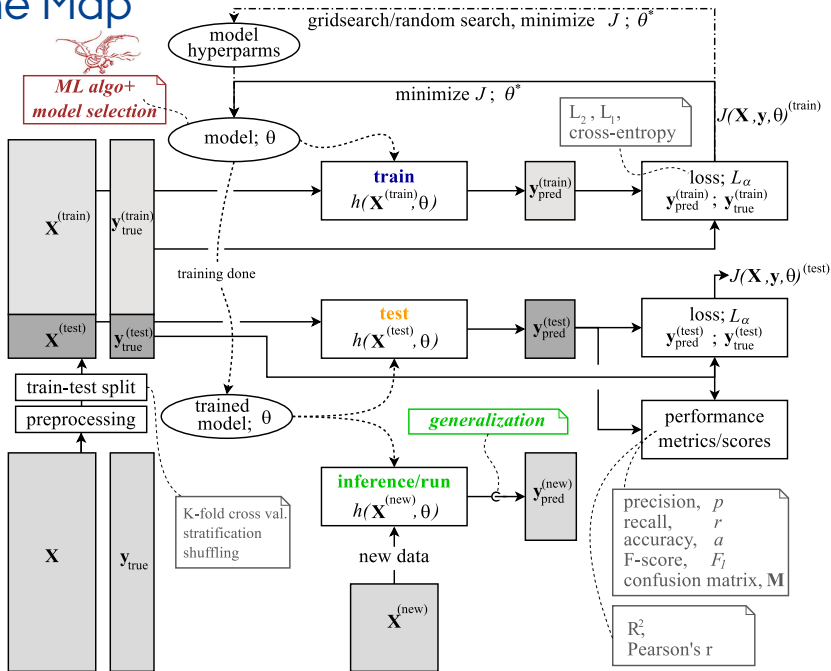
SPRING 2020



# The Map



*ML algo+  
model selection*



# Pipelines

Brief intro to Scikit-learn pipelines..

Python code from capacity\_under\_overfitting.ipynb

```
1  from sklearn.pipeline import Pipeline
2  from sklearn.preprocessing import PolynomialFeatures
3  from sklearn.linear_model import LinearRegression
4  from sklearn.model_selection import cross_val_score
5
6  ..
7
8  polynomial_features = PolynomialFeatures(degree=degrees[i], ..
9  linear_regression = LinearRegression()
10
11  pipeline = Pipeline([
12      ("polynomial_features", polynomial_features),
13      ("linear_regression", linear_regression)
14  ])
15  pipeline.fit(X[:, np.newaxis], y)
16
17  scores = cross_val_score(pipeline, X[:, np.newaxis], y, scoring=
18      "neg_mean_squared_error", cv=10)
19  ..
20  score_mean = -scores.mean()
```

# RESUMÉ: L02/performance\_metrics.ipynb

## Classification metrics

See the [Classification metrics](#) section of the user guide for further details.

|   |  |
|---|--|
| <code>metrics.accuracy_score</code> ( <i>y_true</i> , <i>y_pred</i> [, ...])            | Accuracy classification score.   |
| <code>metrics.auc</code> ( <i>x</i> , <i>y</i> [, <i>reorder</i> ])                     | Compute Area Under the Curve (AUC) using the trapezoidal rule                                    |
| <code>metrics.average_precision_score</code> ( <i>y_true</i> , <i>y_score</i> )         | Compute average precision (AP) from prediction scores  |
| <code>metrics.cohen_kappa_score</code> ( <i>y1</i> , <i>y2</i> [, <i>labels</i> , ...]) | Cohen's kappa: a statistic that measures inter-annotator agreement.                              |
| <code>metrics.confusion_matrix</code> ( <i>y_true</i> , <i>y_pred</i> [, ...])          | Compute confusion matrix to evaluate the accuracy of a classification                            |
| <code>metrics.f1_score</code> ( <i>y_true</i> , <i>y_pred</i> [, <i>labels</i> , ...])  | Compute the F1 score, also known as balanced F-score or F-measure                                |
| <code>metrics.log_loss</code> ( <i>y_true</i> , <i>y_pred</i> [, <i>eps</i> , ...])     | Log loss, aka logistic loss or cross-entropy loss.   |
| <code>metrics.precision_score</code> ( <i>y_true</i> , <i>y_pred</i> [, ...])           | Compute the precision  |
| <code>metrics.recall_score</code> ( <i>y_true</i> , <i>y_pred</i> [, ...])              | Compute the recall   |
| <code>metrics.roc_auc_score</code> ( <i>y_true</i> , <i>y_score</i> [, ...])            | Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores. |
| <code>metrics.roc_curve</code> ( <i>y_true</i> , <i>y_score</i> [, ...])                | Compute Receiver operating characteristic (ROC)  |
| <code>metrics.zero_one_loss</code> ( <i>y_true</i> , <i>y_pred</i> [, ...])             | Zero-one classification loss.  |

## Regression metrics

See the [Regression metrics](#) section of the user guide for further details.



|  |   |
|--|---|
| <code>metrics.explained_variance_score</code> ( <i>y_true</i> , <i>y_pred</i> )  | Explained variance regression score function            |
| <code>metrics.max_error</code> ( <i>y_true</i> , <i>y_pred</i> )                 | max_error metric calculates the maximum residual error. |
| <code>metrics.mean_absolute_error</code> ( <i>y_true</i> , <i>y_pred</i> )       | Mean absolute error regression loss                     |
| <code>metrics.mean_squared_error</code> ( <i>y_true</i> , <i>y_pred</i> [, ...]) | Mean squared error regression loss                      |
| <code>metrics.mean_squared_log_error</code> ( <i>y_true</i> , <i>y_pred</i> )    | Mean squared logarithmic error regression loss          |
| <code>metrics.median_absolute_error</code> ( <i>y_true</i> , <i>y_pred</i> )     | Median absolute error regression loss                   |
| <code>metrics.r2_score</code> ( <i>y_true</i> , <i>y_pred</i> [, ...])           | R^2 (coefficient of determination) regression           |

## Notes on Keras MLPs

Typical Keras MLP Supervised Classifier setup.

- ▶ loss function  
`loss='categorical_crossentropy'`
- ▶ metrics collected via history  
`metrics=['categorical_accuracy', 'mean_squared_error', 'mean_absolute_error']`
- ▶ input lay.: categorical encoding.
- ▶ output lay.: softmax function.

And notice that Keras do *not* provide metrics like precision, recall, F1 but instead `categorical_accuracy`, `binary_accuracy`



# Model capacity

Exercise: `capacity_under_overfitting.ipynb`

Dummy and Paradox classifier:

*capacity* fixed  $\sim 0$ , cannot generalize at all!

Linear regression for a polynomial model:

*capacity*  $\sim$  degree of the polynomial,  $x^n$

Neural Network model:

*capacity*  $\propto$  number of neurons/layers


$\Rightarrow$  **Capacity** can be hard to express as a quantity for some models, but you need to choose..

$\Rightarrow$  how to choose the **optimal capacity**?

# Under- and overfitting

Exercise: `capacity_under_overfitting.ipynb`

Polynomial linear reg. fit for underlying model:  $\cos(x)$



- ▶ underfitting: capacity of model too low,
- ▶ overfitting: capacity too high.


⇒ how to choose the **optimal** capacity?

# Generalization Error


Exercise: `generalization_error.ipynb`

RMSE-capacity plot for lin. reg. with polynomial features

(capacity  $\sim$  degree of poly)



(Figure 5.3 from [DL])



Inspecting the plots from the exercise (`.ipynb`) and [DL],  
extracting the concepts:

- ▶ training/generalization error,
- ▶ generalization gap,
- ▶ underfit/overfit zone,
- ▶ optimal capacity (best-model, early stop),
- ▶ (and the two axes: x/capacity, y/error.)

# Generalization Error

Exercise: `generalization_error.ipynb`

**NOTE:** three methods/plots:

- i) via **learning curves** as in [HOML],
- ii) via an **error-capacity** plot as in [GITHOML] and [DL],
- ii) via an **error-epoch** plot as in [GITHOML].

